



# Bank Deposit Subscription Prediction Report

---

**Team Smart Banker** (Data Science Specialization)

April 26, 2023

# Agenda

---

**01. Meet the Team**

---

**02. Problem Description &  
Business Understanding**

**03 Data Types and  
Problems in the Data**

---

**04. Approach to  
Eliminate Problems in  
the Data**

---

**05. Recommended ML  
Models for the Business  
Problem**

---

**06. Predictions, Results  
and Evaluation**

**06. Reflection and  
Conclusion**

**07. Project Timeline**

# Meet the Team

---

**Group Name:** Team Smart Banker

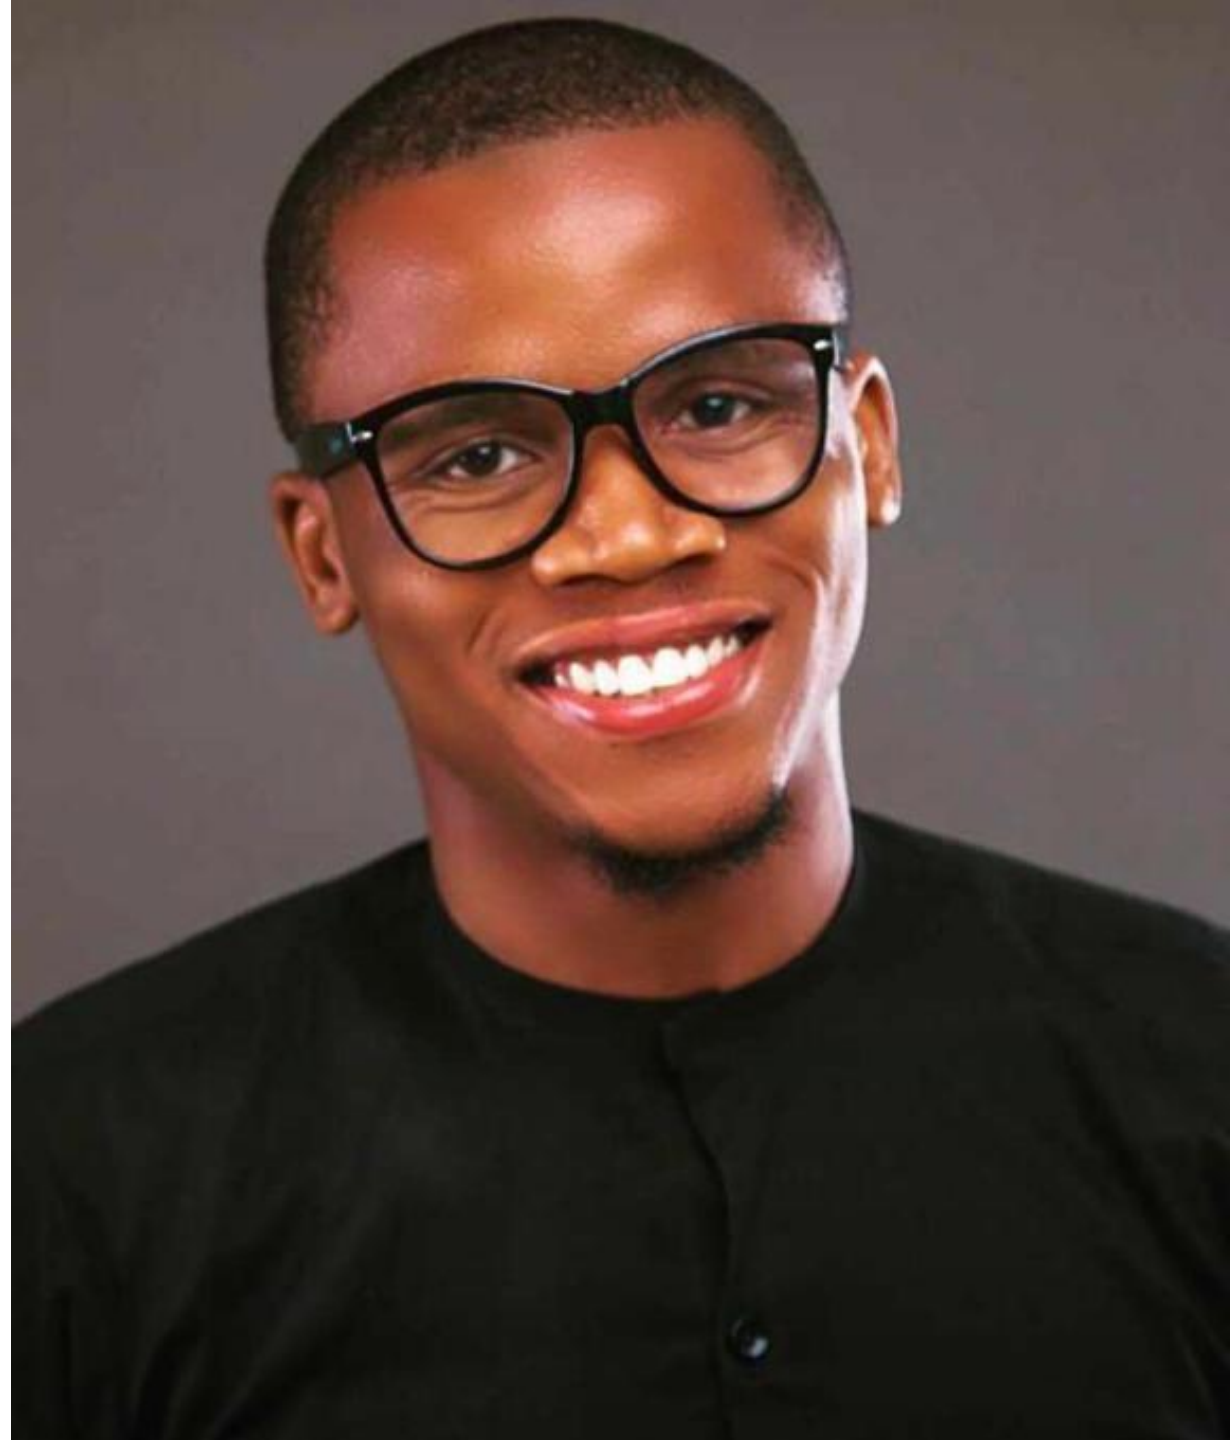
**Member Name:** Chukwujekwu Joseph Ezema

**Email:** [chukwujekwujeks@gmail.com](mailto:chukwujekwujeks@gmail.com)

**School:** Teesside University Middlesbrough

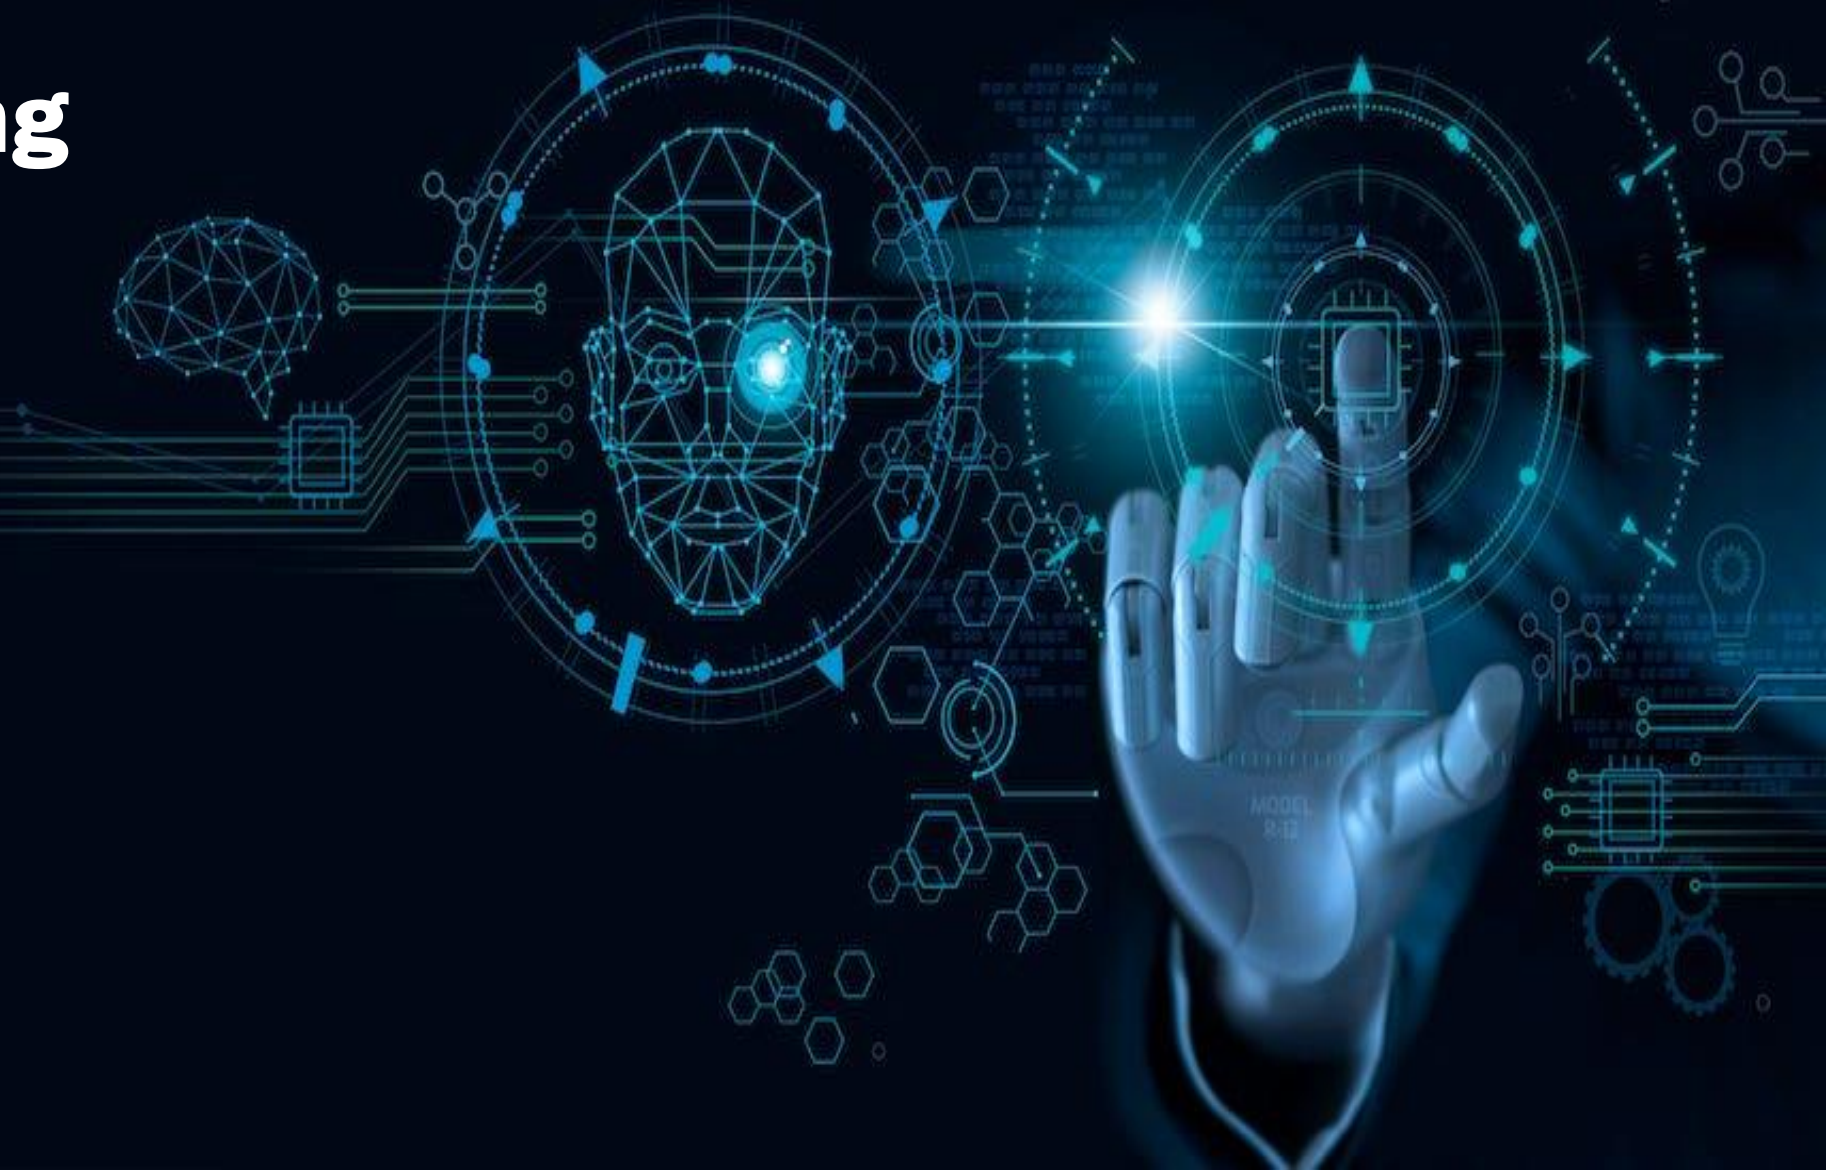
**Country:** United Kingdom

**Specialization:** Data Science



# Bank Marketing Campaign

A Data Science Project



# Problem Description

Objective: Predict if a client will subscribe to a term deposit offered by a bank.

Dataset: Information related to direct marketing campaigns conducted through phone calls by a Portuguese banking institution.

Goal: Build a classification model that predicts whether the client will subscribe to the term deposit or not, based on various input variables such as age, job, marital status, education, etc

# Business Understanding

Importance: Success of direct marketing campaigns depends on accuracy of targeting potential customers who are more likely to purchase the offered product..

Interest: This is to help a bank identify customers who are more likely to subscribe to a term deposit to optimize their marketing strategy and increase their chances of success while minimizing their costs..

# Data Types and Problems in the Data

**Data Type:** The data was majorly dominated by Categorical Data types – which led to most of the them being converted to viable numeric types for ease of modelling for the Classification problem..

## Problem:

1. White spaces in the columns were replaced by “;” thereby making the original data to contain one column instead of 21 columns
2. Some columns were not in their appropriate data types e.g. duration, campaign etc
3. Data contained 0 missing values but 12 duplicated records
4. There were outliers across the numerical columns
5. There is imbalance in the target class, y

```
2]: """ Step 1: Get dataset """  
# 1.2 Load dataset  
  
# Reading the Bank Additional File - dataset for the features  
bank_df = pd.read_csv('bank-additional-full.csv')  
  
bank_df_copy = bank_df.copy() #making a copy  
bank_df_copy
```

```
2]:      age;"job";"marital";"education";"default";"housing";"loan";"contact";"month";"day_of_week";"duration";"campaign";"pdays";"previous";"poutcome";"emp
```

0

1

2

3

4

...

41183

41184

41185

41186

41187

41188 rows × 1 columns

◀

# Original Data Load



In [4]: *# Explore the features of dataset*

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 41188 entries, 0 to 41187
```

```
Data columns (total 21 columns):
```

#	Column	Non-Null Count	Dtype
0	age	41188 non-null	object
1	job	41188 non-null	object
2	marital	41188 non-null	object
3	education	41188 non-null	object
4	default	41188 non-null	object
5	housing	41188 non-null	object
6	loan	41188 non-null	object
7	contact	41188 non-null	object
8	month	41188 non-null	object
9	day_of_week	41188 non-null	object
10	duration	41188 non-null	object
11	campaign	41188 non-null	object
12	pdays	41188 non-null	object
13	previous	41188 non-null	object
14	poutcome	41188 non-null	object
15	emp_var_rate	41188 non-null	object
16	cons_price_idx	41188 non-null	object
17	cons_conf_idx	41188 non-null	object
18	euribor3m	41188 non-null	object
19	nr_employed	41188 non-null	object
20	y	41188 non-null	object

```
dtypes: object(21)
```

```
memory usage: 6.6+ MB
```

## Inappropriate Datatypes

```
In [5]: # Check features with missing value
```

```
data.isnull().sum()
```

```
Out[5]: age          0  
job            0  
marital        0  
education      0  
default        0  
housing        0  
loan           0  
contact        0  
month          0  
day_of_week    0  
duration       0  
campaign       0  
pdays         0  
previous       0  
poutcome       0  
emp_var_rate   0  
cons_price_idx 0  
cons_conf_idx  0  
euribor3m      0  
nr_employed    0  
y              0  
dtype: int64
```

```
In [6]: # Check duplicated records:
```

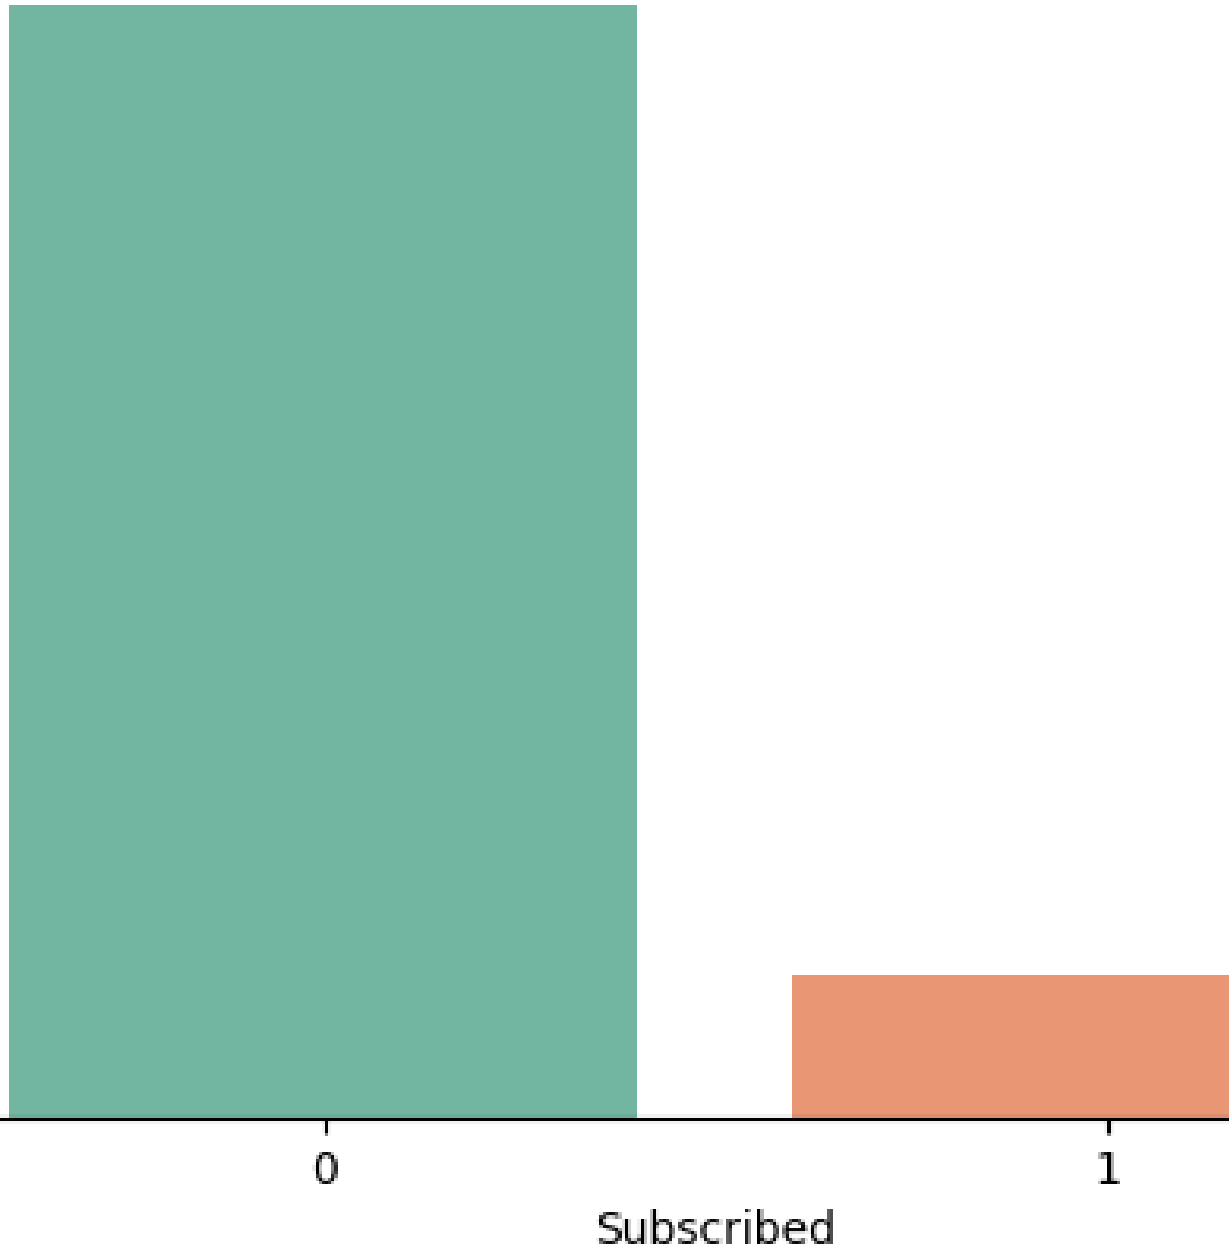
```
print("Number of duplicated records before dropping:
```

```
Number of duplicated records before dropping: 12
```


# Missing & Duplicated Records

---


Unsubscribed Vs Subscribed Customers



**Imbalanced Target Class**



# Outliers



# Approach to Eliminate Problems in the Data

1. Split the columns using the delimiter, semicolon (;) and removed the double the quotes (“”)
2. Use methods from Pandas library to convert to appropriate datatypes e.g `pd.to_numeric()`

3. Dropped duplicated rows
4. Use `StandardScaler()` from sklearn package to limit the effects of the outliers
5. Use `smote()` oversampling technique from imblearn package to handle the imbalance so as not to lose possible data for the training and testing

# Cleaned Dataset

```
In [3]: """ Step 2: Clean dataset """
# Split the columns into separate columns
data = bank_df.iloc[:,0].str.split(";", expand=True)

# Remove the double quotes from the values
data = data.applymap(lambda x: x.strip('\"'))

# Rename the columns with more descriptive names
new_cols = ["age", "job", "marital", "education", "default", "housing", "loan", "contact", "month", "day_of_week", "duration",
            "campaign", "pdays", "previous", "poutcome"]
data = data.rename(columns=dict(zip(data.columns, new_cols)))
data
```

```
Out[3]:
```

	age	job	marital	education	default	housing	loan	contact	month	day_of_week	duration	campaign	pdays	previous	poutcome
0	56	housemaid	married	basic.4y	no	no	no	telephone	may	mon	261	1	999	0	nonexistent
1	57	services	married	high.school	unknown	no	no	telephone	may	mon	149	1	999	0	nonexistent
2	37	services	married	high.school	no	yes	no	telephone	may	mon	226	1	999	0	nonexistent
3	40	admin.	married	basic.6y	no	no	no	telephone	may	mon	151	1	999	0	nonexistent
4	56	services	married	high.school	no	no	yes	telephone	may	mon	307	1	999	0	nonexistent
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
41183	73	retired	married	professional.course	no	yes	no	cellular	nov	fri	334	1	999	0	nonexistent
41184	46	blue-collar	married	professional.course	no	no	no	cellular	nov	fri	383	1	999	0	nonexistent
41185	56	retired	married	university.degree	no	yes	no	cellular	nov	fri	189	2	999	0	nonexistent
41186	44	technician	married	professional.course	no	no	no	cellular	nov	fri	442	1	999	0	nonexistent
41187	74	retired	married	professional.course	no	yes	no	cellular	nov	fri	239	3	999	1	failure

41188 rows × 16 columns

```
# Confirming non duplicated
print("Number of duplicated records after dropping: {}".format(df1.duplicated().sum()))
```

Number of duplicated records after dropping: 0

```
In [8]: """ Step 3: Transform data """
# Convert categorical data to numeric values
def to_integers(self):

    # Converting to binary values for best results
    self.replace({'y' : {'yes' : 1, 'no' : 0}}, inplace=True)

    # Converting to numeric values
    self['age'] = pd.to_numeric(self['age'], errors='coerce')
    self['duration'] = pd.to_numeric(self['duration'], errors='coerce')
    self['campaign'] = pd.to_numeric(self['campaign'], errors='coerce')
    self['pdays'] = pd.to_numeric(self['pdays'], errors='coerce')
    self['previous'] = pd.to_numeric(self['previous'], errors='coerce')
    self['emp_var_rate'] = pd.to_numeric(self['emp_var_rate'], errors='coerce')
    self['cons_price_idx'] = pd.to_numeric(self['cons_price_idx'], errors='coerce')
    self['cons_conf_idx'] = pd.to_numeric(self['cons_conf_idx'], errors='coerce')
    self['euribor3m'] = pd.to_numeric(self['euribor3m'], errors='coerce')
    self['nr_employed'] = pd.to_numeric(self['nr_employed'], errors='coerce')

    return self

to_integers(df1)
```

## Dropped Duplicates and converted to numeric

Out[8]:

	age	job	marital	education	default	housing	loan	contact	month	day_of_week	duration	campaign	pdays	previous	outcome
0	56	housemaid	married	basic.4y	no	no	no	telephone	may	mon	261	1	999	0	nonexistent
1	57	services	married	high.school	unknown	no	no	telephone	may	mon	149	1	999	0	nonexistent
2	37	services	married	high.school	no	yes	no	telephone	may	mon	226	1	999	0	nonexistent
3	40	admin	married	basic.6y	no	no	no	telephone	may	mon	151	1	999	0	nonexistent

# After Transformation

- Cleaned the jam-packed columns into distinct columns
- Removed duplicates
- Converted wrong object datatypes to appropriate numerical types
- Created a function to label encode and further applied one-hot encoding using `pd.get_dummies` to put the categories into 1's and 0's in order to avoid biases
- Note: The increased number of columns will be tackled with dimensionality reduction strategy like PCA

In [11]: *# confirm the nature of the data after feature engineering*

```
df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

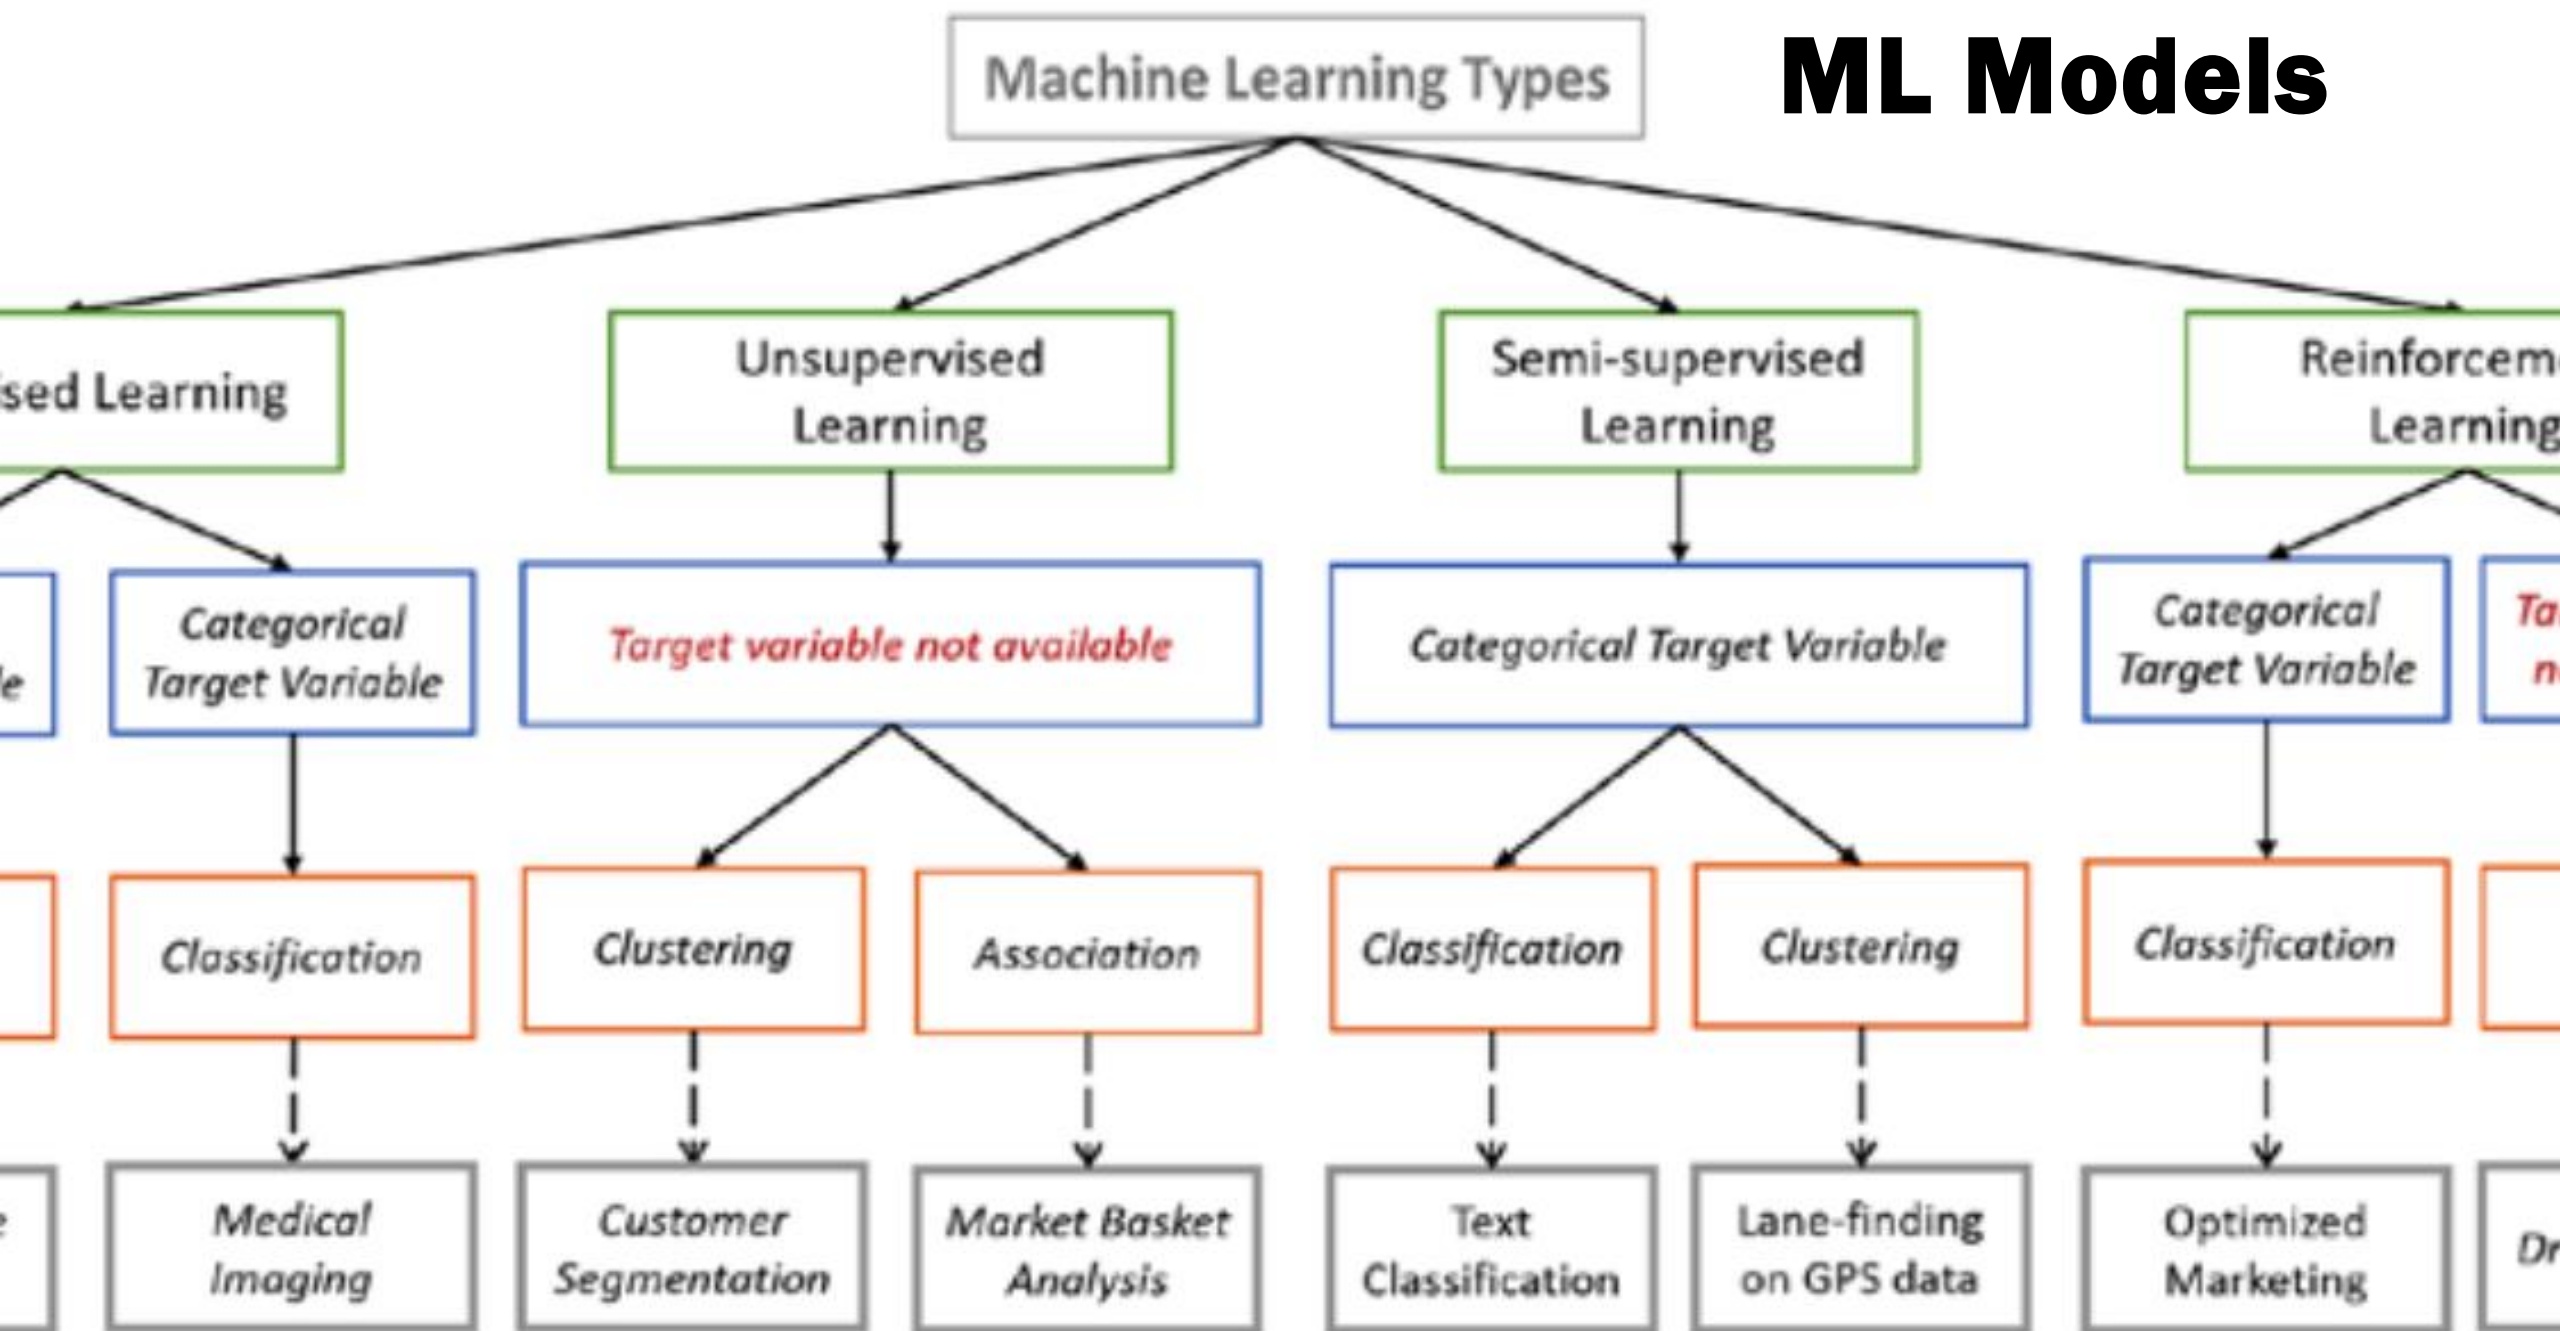
```
Int64Index: 41176 entries, 0 to 41187
```

```
Data columns (total 44 columns):
```

#	Column	Non-Null Count	Dtype
0	Age	41176 non-null	int64
1	Duration	41176 non-null	int64
2	Campaign	41176 non-null	int64
3	Pdays	41176 non-null	int64
4	Previous	41176 non-null	int64
5	Emp_var_rate	41176 non-null	float64
6	Cons_price_idx	41176 non-null	float64
7	Cons_conf_idx	41176 non-null	float64
8	Euribor3m	41176 non-null	float64
9	No_employed	41176 non-null	float64
10	Subscribed	41176 non-null	int64
11	Job_Employed	41176 non-null	uint8
12	Job_Self_employed	41176 non-null	uint8
13	Job_Unemployed	41176 non-null	uint8
14	Job_Unknown	41176 non-null	uint8
15	Marital_Married	41176 non-null	uint8
16	Marital_Not_married	41176 non-null	uint8
17	Marital_Unknown	41176 non-null	uint8
18	Education_Basic	41176 non-null	uint8
19	Education_Secondary	41176 non-null	uint8
20	Education_Tertiary	41176 non-null	uint8
21	Education_Uneducated	41176 non-null	uint8
22	Education_Unknown	41176 non-null	uint8
23	Default_no	41176 non-null	uint8
24	Default_unknown	41176 non-null	uint8
25	Default_yes	41176 non-null	uint8
26	Housing_no	41176 non-null	uint8
27	Housing_unknown	41176 non-null	uint8
28	Housing_yes	41176 non-null	uint8
29	Loan_no	41176 non-null	uint8
30	Loan_unknown	41176 non-null	uint8



# ML Models



# ML Model Recommendations

1. Logistics Regression – Linear
2. K-Nearest Neighbour Classifier – Non-parametric
3. Random Forest Classifier – Bagging Ensemble

4. Gradient Boost Classifier – Boosting Ensemble
5. XGBoost Classifier – Boosting Ensemble
6. Neural Networks (Keras) for Deep Learning – Non-linear

# Why Choose them?

```
""" Step 6: Model the Algorithms for Prediction """  
  
# Design ML model - Set up  
  
# Classifiers  
classifiers = {  
    "LogisticRegression": LogisticRegression(random_state=random_state),  
    "KNearest": KNeighborsClassifier(),  
    "GradientBoost": GradientBoostingClassifier(random_state=random_state),  
    "Random Forest Classifier": RandomForestClassifier(random_state=random_state),  
    "XGBClassifier": XGBClassifier(random_state=random_state)  
}
```

- **Logistics Regression:** Linear model. Logistic regression is commonly used in classification problems because it provides a probabilistic interpretation of the outputs, allowing for easy thresholding for binary classification tasks. It can also handle imbalanced classes by adjusting the threshold for classification.
- **Gradient Boost Classifier:** Ensemble model. Gradient boosting is a technique that combines multiple weak learners to create a strong learner. It can handle imbalanced data because it can assign more weight to the minority class during the boosting process
- **Neural Networks (Keras) for Deep Learning:** Non-linear model. Neural networks are powerful models that can learn complex patterns in data. They can handle imbalanced data through techniques such as class weighting, oversampling of the minority class, and the use of specialized loss functions

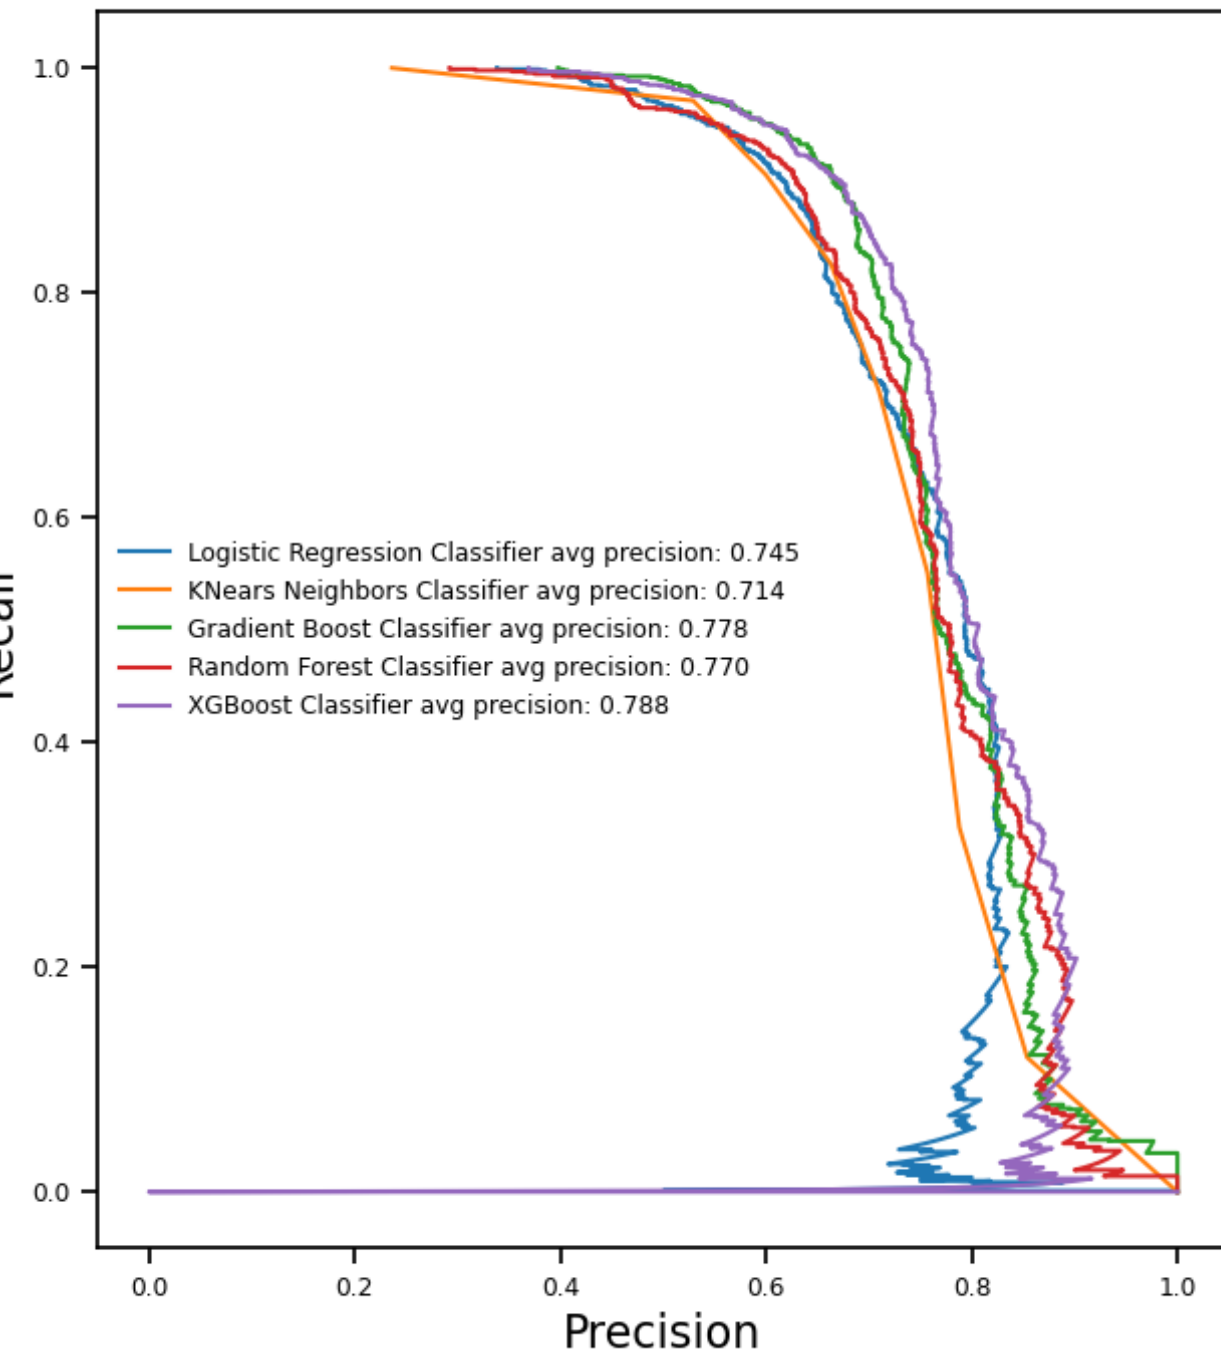
```
# Setup for Deep Learning Model  
def ANN_model(X_train, y_train, X_test, y_test, epochs=5):  
  
    # design model  
    model = keras.Sequential([  
        keras.layers.Dense(18, input_shape=(15,), activation='relu'),  
        keras.layers.Dense(1, activation='sigmoid')  
    ])  
  
    # compile model  
    model.compile(optimizer = 'adam',  
                  loss = 'binary_crossentropy',  
                  metrics = ['accuracy'])  
  
    # fit model  
    model.fit(X_train, y_train, epochs=epochs)  
  
    # evaluate model  
    print(f'model evaluation :', model.evaluate(X_test, y_test))  
  
    # build prediction series  
    yp = model.predict(X_test)  
    y_pred = []  
    for element in yp:  
        if element > 0.5:  
            y_pred.append(1)  
        else:  
            y_pred.append(0)  
  
    mse = np.mean(np.power(X_test - yp, 2), axis=1)  
    error = pd.DataFrame({'reconstruction_error': mse,  
                          'true_class': y_test})  
  
    print(error.describe())  
  
    # result  
    print(classification_report(y_test, y_pred, labels=[1,0]))
```

# Predictions, Results and Evaluation

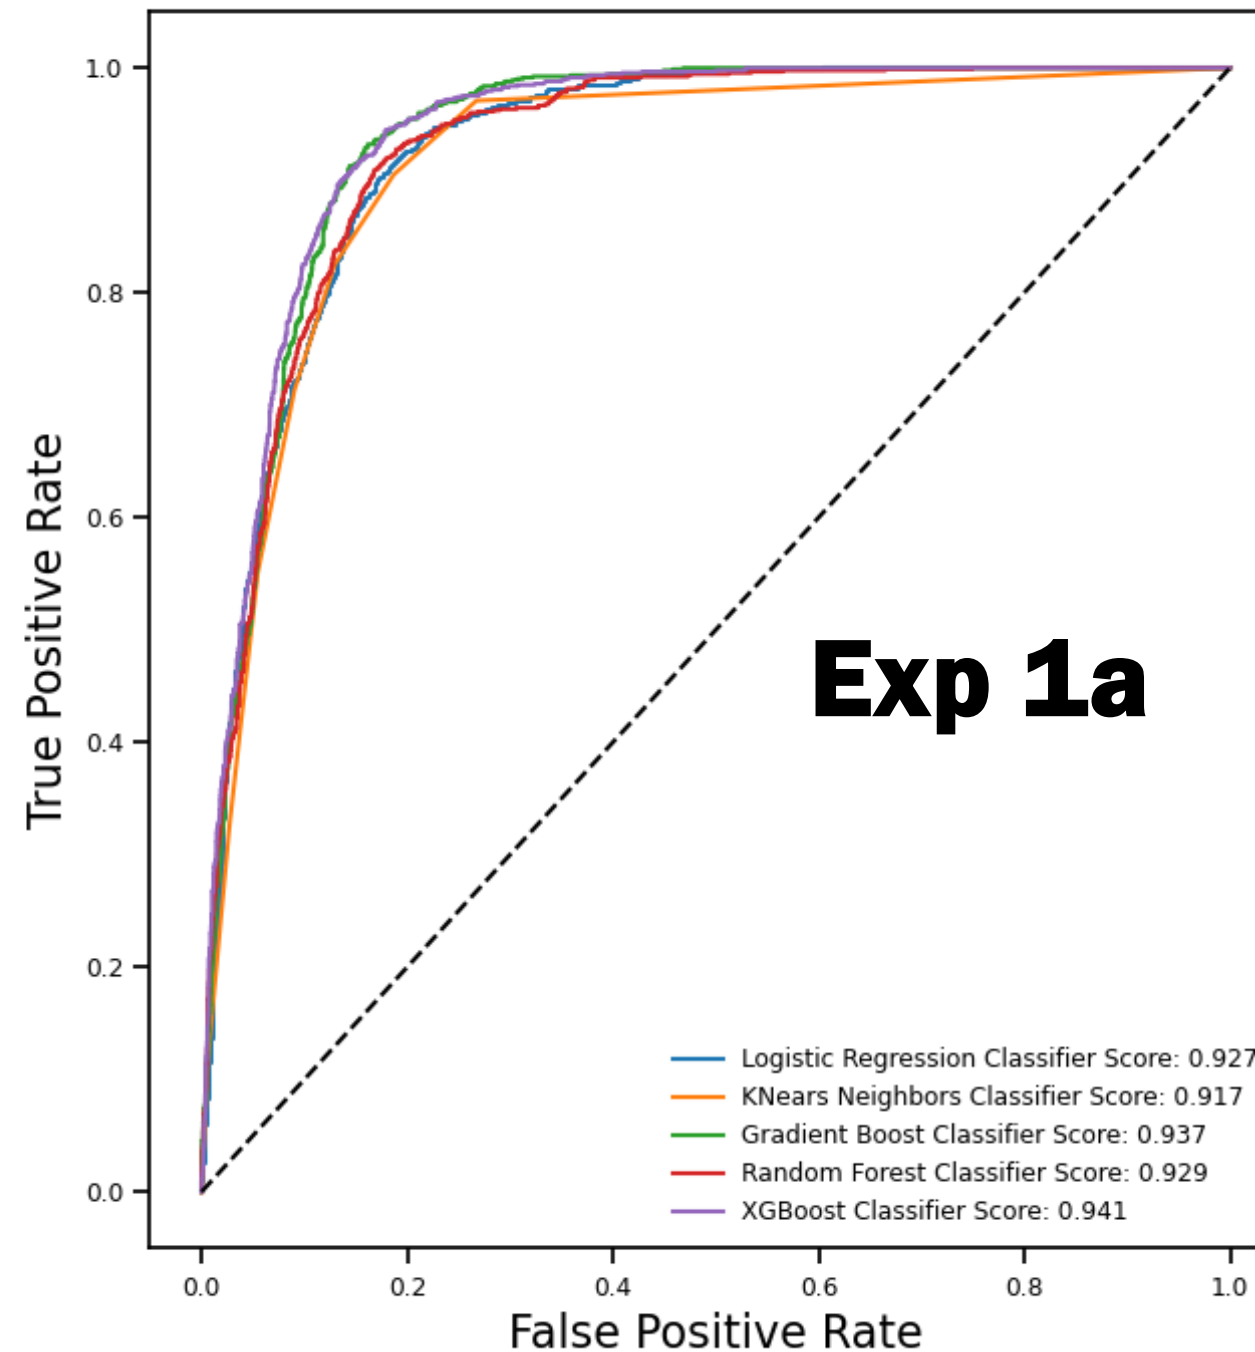
Exp 1a: Classifiers on Best Features of Normal Sample Dataset (without oversampling)

Classifier	Accur. Score	Macro F1- Score	Weighted F1-Score	P-R Curve	AUC- ROC
Log. Regression	0.87	0.80	0.86	0.75	0.93
KNN	0.87	0.82	0.87	0.71	0.92
Gradient Boost	0.89	0.85	0.89	0.78	0.94
Random Forest	0.86	0.79	0.85	0.77	0.93
XG Boost	0.89	0.85	0.89	0.79	0.94

Precision-Recall Curve



ROC Curve




# Predictions, Results and Evaluation


Exp 1b: PCA applied on Normal Sample Dataset (without oversampling)

Classifier	Accur. Score	Macro F1- Score	Weighted F1-Score	P-R Curve	AUC- ROC
Log. Regression	0.87	0.81	0.86	0.74	0.92
KNN	0.86	0.80	0.85	0.68	0.89
Gradient Boost	0.87	0.82	0.87	0.75	0.92
Random Forest	0.85	0.76	0.84	0.73	0.91
XG Boost	0.88	0.83	0.88	0.74	0.92

Precision-Recall Curve



ROC Curve



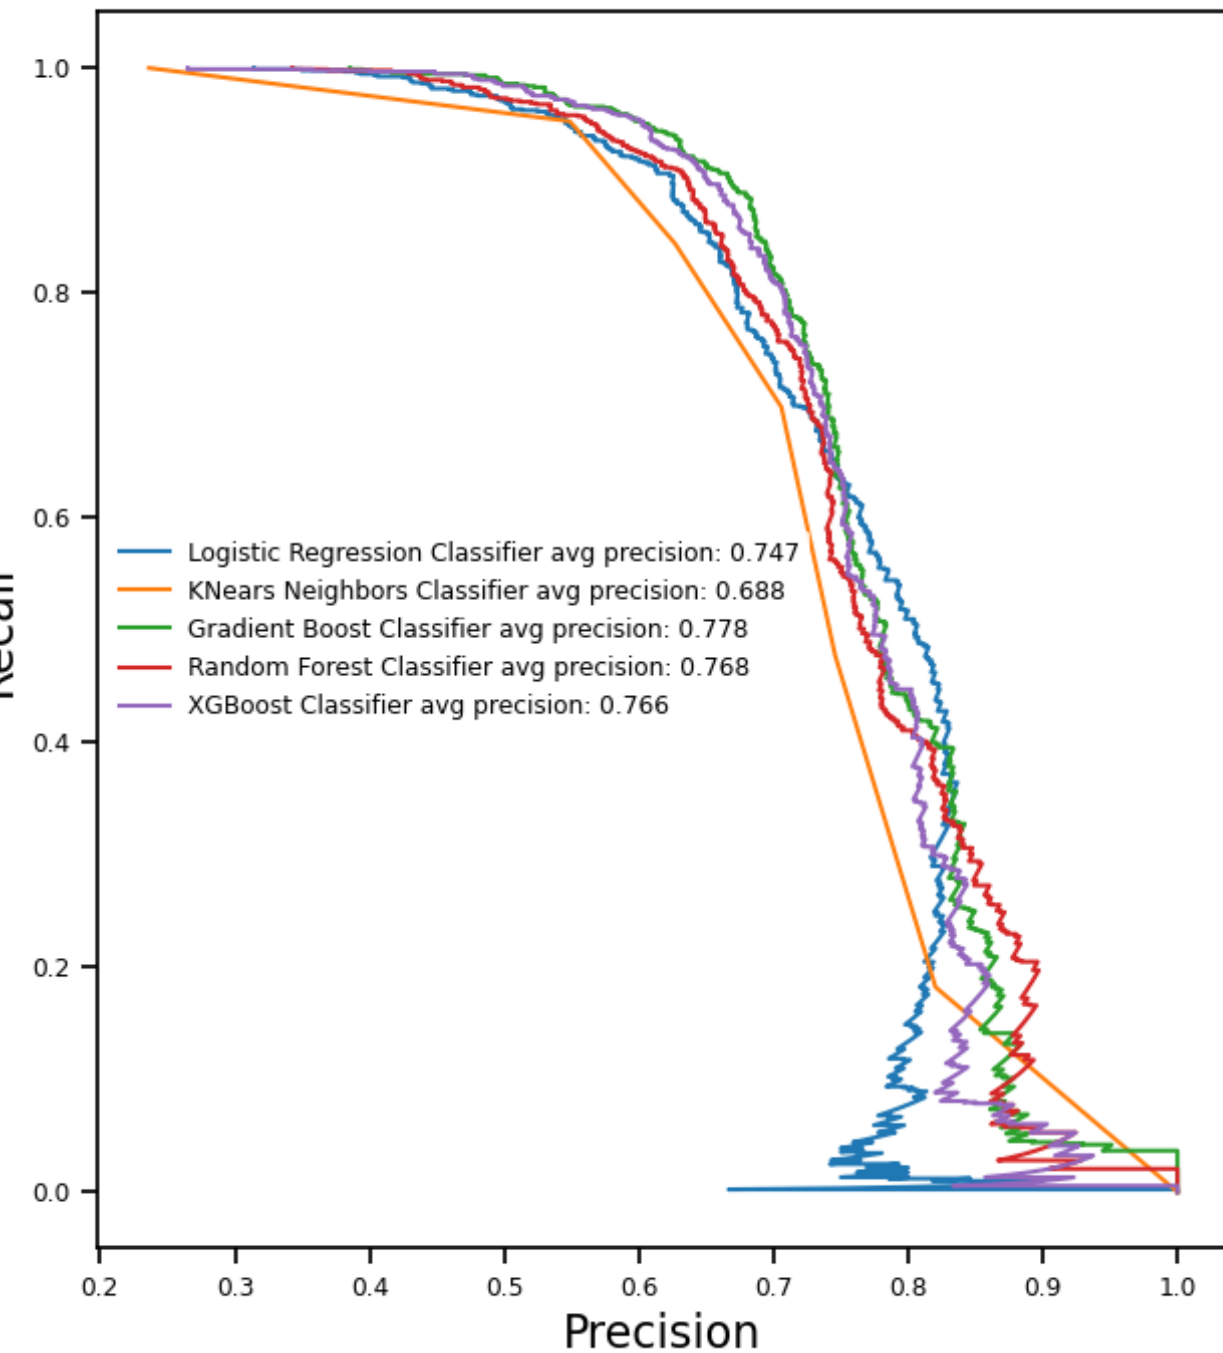
# Predictions, Results and Evaluation

Exp 2a: Smote Technique on Best Features - Oversampling


Classifier	Accur. Score	Macro F1- Score	Weighted F1-Score	P-R Curve	AUC- ROC
Log. Regression	0.86	0.82	0.86	0.75	0.93
KNN	0.85	0.82	0.86	0.69	0.91
Gradient Boost	0.88	0.85	0.89	0.78	0.94
Random Forest	0.84	0.81	0.85	0.77	0.93
XG Boost	0.88	0.85	0.89	0.77	0.94



Precision-Recall Curve



ROC Curve




# Predictions, Results and Evaluation


Exp 2b: SMOTE Technique on PCA - Oversampling

Classifier	Accur. Score	Macro F1-Score	Weighted F1-Score	P-R Curve	AUC-ROC
Log. Regression	0.86	0.82	0.86	0.74	0.92
KNN	0.85	0.80	0.85	0.60	0.85
Gradient Boost	0.87	0.83	0.87	0.75	0.93
Random Forest	0.83	0.80	0.84	0.74	0.92
XG Boost	0.87	0.83	0.87	0.74	0.92

Precision-Recall Curve



ROC Curve



# Predictions, Results and Evaluation

Exp 3a/3b: Artificial Neural Network Classification Model

Sample	Accur. Score	Macro F1- Score	Weighted F1-Score	P-R Curve	AUC- ROC
Best Features	0.87	0.82	0.87	0.63	0.81
SMOTE Sample	0.86	0.83	0.87	0.66	0.85

# Reflection and Conclusion

This project was to build a classification model, at least one from a family of algorithm, that predicts whether clients will subscribe to the term deposit or not, based on various input variables such as age, job, marital status, education etc. The dataset was properly cleaned and preprocessed. To reduce training time, a resampling was carried out, this was also to reduce the adverse effect of imbalance in the data. The data were scaled using StandardScaler. The feature selection technique proved to be quite good on the models, with less impact by SMOTE techniques and PCA feature selection.

GridSearchCV was used to get the best-fit parameters on each of the models, which was evident in the model evaluation of 70 to 90% at almost all levels. The results from the different experiments and classification models didn't show much differences except the impacts by ensemble boosting models. In conclusion, the prediction was a success as proper machine learning workflow was followed for a better result. The imbalanced dataset was handled by resampling technique – which can heavily affect model evaluation. Thus, these steps are highly recommended for classification problems.

# Project Timeline

## Wk1. 19 Mar – 26 Mar 2023

**Data Collection:** Download the dataset from the given source and load it into a Jupyter notebook.

**Data Exploration:** Perform preliminary data exploration to understand the dataset's structure, quality, and potential challenges.

## Wk2. 26 Mar – 2 Apr 2023

**Data Preparation:** Prepare the dataset for modeling by performing data cleaning, feature selection, feature engineering, and data transformation as necessary.

**Modeling:** Train different classification models on the prepared dataset and evaluate their performance using appropriate metrics.

## Wk3. 2 Apr – 9 Apr 2023

**Model Tuning:** Tune the hyperparameters of the best-performing model to improve its performance further.

**Model Interpretation:** Interpret the trained model and identify the most significant factors that contribute to the prediction.

## Wk4. 9 Apr – 16 Apr 2023

**Finalize the project report:** Create a comprehensive report documenting the entire project, including problem definition, methodology, results, and interpretation.

**Submission:** Submit the final report within 2 weeks by April 28, 2023, as per the given deadline..



# Thank you

---

Thanks for going through this interesting journey with us.  
We hope we have added value to this work.

GitHub Link

<https://github.com/Jeks042/Data-Science-Projects>